

## Lesson 03 Demo 03

### Launching a Pod and Establishing an Associated Service

**Objective:** To demonstrate how to seamlessly integrate Kubernetes deployments with services to achieve scalable and accessible pod configurations

**Tools required:** kubeadm, kubectl, kubelet, and containerd

**Prerequisites:** A Kubernetes cluster should already be set up (refer to the steps provided in Lesson 02, Demo 01 for guidance).

Steps to be followed:

1. Create a deployment object
2. Create a service with label selector for deployment

#### Step 1: Create a deployment object

- 1.1 Create a YAML file for the deployment using the following command:  
**vi mydeployment.yaml**

```
labsuser@master:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
master.example.com                 Ready    control-plane   9d   v1.28.2
worker-node-1.example.com          Ready    <none>         9d   v1.28.2
worker-node-2.example.com          Ready    <none>         9d   v1.28.2
labsuser@master:~$ vi mydeployment.yaml
```

1.2 Add the following code to the **mydeployment.yaml**:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: httpd
  replicas: 2
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
      - name: httpd
        image: httpd:latest
        ports:
        - containerPort: 80
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: httpd
  replicas: 2
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
      - name: httpd
        image: httpd:latest
        ports:
        - containerPort: 80
```

- 1.3 Apply the deployment object using the following command:  
**kubectl apply -f mydeployment.yaml**

```
labsuser@master:~$ vi mydeployment.yaml
labsuser@master:~$ kubectl apply -f mydeployment.yaml
deployment.apps/nginx-deployment created
labsuser@master:~$
```

- 1.4 Verify the deployment and its pods using the following commands:  
**kubectl get deployment**  
**kubectl get pods**

```
labsuser@master:~$ kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
admin         0/1     1            0           5d22h
nginx         1/1     1            1           4d4h
nginx-deployment 2/2     2            2           23s
labsuser@master:~$ kubectl get pods
NAME                                READY   STATUS             RESTARTS   AGE
admin-56d684dff9-zjfhc              0/1     ImagePullBackOff   0          5d22h
counter                             1/1     Running            4 (27m ago) 5d2h
nginx-7854ff8877-mvrtr               1/1     Running            1 (4d ago)  4d4h
nginx-deployment-6d6b866d8f-bw8xr    1/1     Running            0          38s
nginx-deployment-6d6b866d8f-r7pnj    1/1     Running            0          38s
pod-demo                             1/1     Running            8 (27m ago) 9d
labsuser@master:~$
```

## Step 2: Create a service with label selector for deployment

- 2.1 Create a new YAML file for the service using the command below:  
**vi myservice.yaml**

```
labsuser@master:~$ kubectl get pods
NAME                                READY   STATUS             RESTARTS   AGE
admin-56d684dff9-zjfhc              0/1     ImagePullBackOff   0          5d22h
counter                             1/1     Running            4 (27m ago) 5d2h
nginx-7854ff8877-mvrtr               1/1     Running            1 (4d ago)  4d4h
nginx-deployment-6d6b866d8f-bw8xr    1/1     Running            0          38s
nginx-deployment-6d6b866d8f-r7pnj    1/1     Running            0          38s
pod-demo                             1/1     Running            8 (27m ago) 9d
labsuser@master:~$ vi myservice.yaml
```

2.2 Add the following code to the **myservice.yaml**:

```
apiVersion: v1
kind: Service
metadata:
  name: myservice
spec:
  selector:
    app: httpd
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: myservice
spec:
  selector:
    app: httpd
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

2.3 Apply the service object using the following command:

**kubectl apply -f myservice.yaml**

```
labsuser@master:~$ vi myservice.yaml
labsuser@master:~$ kubectl apply -f myservice.yaml
service/myservice created
labsuser@master:~$
```

- 2.4 Describe the service to verify its connection to the pods using the command below:  
**kubectl describe svc myservice**

```
labsuser@master:~$ kubectl apply -f myservice.yaml
service/myservice created
labsuser@master:~$ kubectl describe svc myservice
Name:                myservice
Namespace:            default
Labels:               <none>
Annotations:          <none>
Selector:             app=httpd
Type:                 ClusterIP
IP Family Policy:     SingleStack
IP Families:          IPv4
IP:                   10.99.56.93
IPs:                  10.99.56.93
Port:                 <unset> 8080/TCP
TargetPort:           80/TCP
Endpoints:            192.168.232.208:80,192.168.47.146:80
Session Affinity:     None
Events:               <none>
labsuser@master:~$
```

- 2.5 Check the targeted pods by listing them using the service's selector using the command below:  
**kubectl get pods -l app=httpd**

```
labsuser@master:~$ kubectl get pods -l app=httpd
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6d6b866d8f-bw8xr   1/1     Running   0           4m55s
nginx-deployment-6d6b866d8f-r7pnj   1/1     Running   0           4m55s
labsuser@master:~$
```

2.6 List the service's endpoints to view the IP addresses of the pods it targets:

**kubectl get endpoints myservice**

```
labsuser@master:~$ kubectl get endpoints myservice
NAME           ENDPOINTS                                     AGE
myservice      192.168.232.208:80,192.168.47.146:80       2m39s
labsuser@master:~$
```

By following these steps, you have successfully set up a Kubernetes deployment that offers scalability and a service that ensures the deployed pods accessibility within the cluster or to external clients.