

Lesson 05 Demo 02

Configuring Pod Affinity and Anti-Affinity in Kubernetes

Objective: To configure pod affinity and anti-affinity rules in a Kubernetes cluster to ensure specific deployment patterns of pods across nodes

Tools required: kubeadm, kubectl, kubelet, and containerd

Prerequisites: A Kubernetes cluster should already be set up (refer to the steps in Lesson 02, Demo 01 for guidance).

Steps to be followed:

1. Deploy redis-cache with anti-affinity
2. Colocate web server with redis-cache using affinity

Step 1: Deploy redis-cache with anti-affinity

- 1.1 Create the redis-cache-deployment.yaml configuration file by entering the command:
`vi redis-cache-deployment.yaml`

```
labsuser@master:~$ vi redis-cache-deployment.yaml
labsuser@master:~$
```

1.2 Paste the following code to the **redis-cache-deployment.yaml** file:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-cache
spec:
  selector:
    matchLabels:
      app: store

  replicas: 3
  template:
    metadata:
      labels:
        app: store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - store
              topologyKey: "kubernetes.io/hostname"
      containers:
        - name: redis-server
          image: redis:3.2-alpine
```

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-cache
spec:
  selector:
    matchLabels:
      app: store
  replicas: 3
  template:
    metadata:
      labels:
        app: store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - store
              topologyKey: "kubernetes.io/hostname"
      containers:
        - name: redis-server
          image: redis:3.2-alpine

```

- 1.3 Apply the **redis-cache-deployment.yaml** configuration file using the command:
kubectl apply -f redis-cache-deployment.yaml

```

labsuser@master:~$ kubectl apply -f redis-cache-deployment.yaml
deployment.apps/redis-cache created
labsuser@master:~$

```

- 1.4 Verify the deployment of **redis-cache** with the following commands:
kubectl get deploy redis-cache
kubectl get pod -l app=store -o wide

```

labsuser@master:~$ kubectl apply -f redis-cache-deployment.yaml
deployment.apps/redis-cache created
labsuser@master:~$ kubectl get deploy redis-cache
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
redis-cache 1/3      3            1           6m43s
labsuser@master:~$ kubectl get pod -l app=store -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE   READINESS GATES
redis-cache-8478cbd86-74wrv          0/1     Pending   0           6m56s   <none>        <none>                <none>            <none>
redis-cache-8478cbd86-qw8jb          0/1     Pending   0           6m56s   <none>        <none>                <none>            <none>
redis-cache-8478cbd86-wldjq          1/1     Running   0           6m57s   172.16.232.202 worker-node-2.example.com          <none>            <none>

```

Step 2: Colocate web server with redis-cache using affinity

2.1 Create the **web-server-deployment.yaml** configuration file by using the command:
vi web-server-deployment.yaml

```
labsuser@master:~$ kubectl apply -f redis-cache-deployment.yaml
deployment.apps/redis-cache created
labsuser@master:~$ kubectl get deploy redis-cache
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
redis-cache   1/3     3            1           6m43s
labsuser@master:~$ kubectl get pod -l app=store -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                NOMINATED NODE   READINESS GATES
redis-cache-8478cbdc86-74wnv  0/1     Pending   0           6m56s   <none>        <none>              <none>           <none>
redis-cache-8478cbdc86-qw8jb  0/1     Pending   0           6m56s   <none>        <none>              <none>           <none>
redis-cache-8478cbdc86-wldjq  1/1     Running   0           6m57s   172.16.232.202  worker-node-2.example.com  <none>           <none>
labsuser@master:~$ kubectl get pod -l app=store -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                NOMINATED NODE   READINESS GATES
redis-cache-8478cbdc86-74wnv  0/1     Pending   0           17m   <none>        <none>              <none>           <none>
redis-cache-8478cbdc86-qw8jb  0/1     Pending   0           17m   <none>        <none>              <none>           <none>
redis-cache-8478cbdc86-wldjq  1/1     Running   0           17m   172.16.232.202  worker-node-2.example.com  <none>           <none>
labsuser@master:~$ vi web-server-deployment.yaml
labsuser@master:~$
```

2.2 Paste the following code to the **web-server-deployment.yaml** file:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-server
spec:
  selector:
    matchLabels:
      app: web-store
  replicas: 3
  template:
    metadata:
      labels:
        app: web-store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - web-store
              topologyKey: "kubernetes.io/hostname"
        podAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - store
              topologyKey: "kubernetes.io/hostname"
      containers:
        - name: web-app
          image: nginx:1.16-alpine
```

```

template:
  metadata:
    labels:
      app: web-store
  spec:
    affinity:
      podAntiAffinity:
        requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
                - key: app
                  operator: In
                  values:
                    - web-store
              topologyKey: "kubernetes.io/hostname"
      podAffinity:
        requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
                - key: app
                  operator: In
                  values:
                    - store
              topologyKey: "kubernetes.io/hostname"
    containers:
      - name: web-app
        image: nginx:1.16-alpine

```

:wq

2.3 Apply the **web-server-deployment.yaml** configuration using the command:
kubectl apply -f web-server-deployment.yaml

```

redis-cache-8478cbdc86-qw8jb 0/1 Pending 0 17m <none>
redis-cache-8478cbdc86-wldjq 1/1 Running 0 17m 172.16.232.202
labsuser@master:~$ vi web-server-deployment.yaml
labsuser@master:~$ kubectl apply -f web-server-deployment.yaml
deployment.apps/web-server created
labsuser@master:~$

```

2.4 Verify the deployment of the web server with the following commands:

```
kubectl get deploy web-server
kubectl get pod -l app=web-store -o wide
```

```
labsuser@master:~$ kubectl get deploy web-server
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
web-server 1/3      3            1           95s
labsuser@master:~$ kubectl get pod -l app=web-store -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE   READINESS GATES
web-server-55f57c89d4-8l1nb         0/1     Pending   0           103s   <none>        <none>                                <none>            <none>
web-server-55f57c89d4-kh5st         1/1     Running   0           103s   172.16.232.203 worker-node-2.example.com          <none>            <none>
web-server-55f57c89d4-rbxrf         0/1     Pending   0           103s   <none>        <none>                                <none>            <none>
```

2.5 Check the information of the pods using the following commands:

```
kubectl get pods -l app=store -o wide
kubectl get pods -l app=web-store -o wide
```

```
labsuser@master:~$ kubectl get deploy web-server
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
web-server 1/3      3            1           95s
labsuser@master:~$ kubectl get pod -l app=web-store -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE   READINESS GATES
web-server-55f57c89d4-8l1nb         0/1     Pending   0           103s   <none>        <none>                                <none>            <none>
web-server-55f57c89d4-kh5st         1/1     Running   0           103s   172.16.232.203 worker-node-2.example.com          <none>            <none>
web-server-55f57c89d4-rbxrf         0/1     Pending   0           103s   <none>        <none>                                <none>            <none>
labsuser@master:~$ kubectl get pods -l app=store -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE   READINESS GATES
redis-cache-8478cbdc86-74wv         0/1     Pending   0           27m   <none>        <none>                                <none>            <none>
redis-cache-8478cbdc86-qw8jb         0/1     Pending   0           27m   <none>        <none>                                <none>            <none>
redis-cache-8478cbdc86-wldjq         1/1     Running   0           27m   172.16.232.202 worker-node-2.example.com          <none>            <none>
labsuser@master:~$ kubectl get pods -l app=web-store -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE   READINESS GATES
web-server-55f57c89d4-8l1nb         0/1     Pending   0           2m48s <none>        <none>                                <none>            <none>
web-server-55f57c89d4-kh5st         1/1     Running   0           2m48s 172.16.232.203 worker-node-2.example.com          <none>            <none>
web-server-55f57c89d4-rbxrf         0/1     Pending   0           2m48s <none>        <none>                                <none>            <none>
```

By following these steps, you have successfully configured pod affinity and anti-affinity in Kubernetes, ensuring that your **redis-cache** instances are spread across different hosts and your **web-server** instances are colocated with the **redis-cache** in the same nodes for optimal performance and resilience.